

# Define a new and remove a view

```
CREATE VIEW venue_view_2017 AS  
  SELECT *  
  FROM venue  
  WHERE year = 2017;
```

```
DROP VIEW venue_view_2017;
```

For more information:

- <https://www.postgresql.org/docs/current/static/sql-createview.html>
- <https://www.postgresql.org/docs/current/static/sql-dropview.html>

# View

The view is not physically materialized. Instead, the query is run every time the view is referenced in a query.

```
SELECT * FROM venue_view_2017;
```

id	name	year	school	volume	number	type
(0 rows)						

```
INSERT INTO venue VALUES (3150859, 'VLDB', 2017, '', '9', '12', 0);
```

```
SELECT * FROM venue_view_2017;
```

id	name	year	school	volume	number	type
3150859	VLDB	2017		9	12	0
(1 row)						

# Materialized view

The query is executed and used to populate the view at the time the command is issued and may be refreshed later using `REFRESH MATERIALIZED VIEW`.

`CREATE MATERIALIZED VIEW` is similar to `CREATE TABLE AS`, except that it also remembers the query used to initialize the view, so that it can be refreshed later upon demand.

A materialized view has many of the same properties as a table, but there is no support for temporary materialized views or automatic generation of OIDs.

# Define a new, remove and refresh a materialized view

```
CREATE MATERIALIZED VIEW coauthors_per_paper AS  
  SELECT paperid AS paper, COUNT(authid) AS coauthors  
  FROM paperauths  
  GROUP BY paperid;
```

```
REFRESH MATERIALIZED VIEW coauthors_per_paper;
```

```
DROP MATERIALIZED VIEW coauthors_per_paper;
```

For more information:

- <https://www.postgresql.org/docs/current/static/sql-creatematerializedview.html>
- <https://www.postgresql.org/docs/current/static/sql-refreshmaterializedview.html>
- <https://www.postgresql.org/docs/current/static/sql-dropmaterializedview.html>

# Materialized view

```
SELECT * FROM coauthors_per_paper WHERE paper = 342;
```

paper	coauthors
342	2

(1 row)

```
INSERT INTO paperauths VALUES (342, 43);
```

```
SELECT * FROM coauthors_per_paper WHERE paper = 342;
```

paper	coauthors
342	2

(1 row)

# Materialized view (cont)

```
REFRESH MATERIALIZED VIEW coauthors_per_paper;
```

```
REFRESH MATERIALIZED VIEW
```

```
SELECT * FROM coauthors_per_paper WHERE paper = 342;
```

paper	coauthors
342	3

(1 row)


# Define a new table as

**CREATE TABLE AS** creates a table and fills it with data computed by a **SELECT** command.

The table columns have the names and data types associated with the output columns of the **SELECT**.

```
CREATE TABLE journal_articles AS  
  SELECT p.*  
  FROM papers p, venue v  
  WHERE p.venue = v.id AND v.type = 0;
```

For more information:

 <https://www.postgresql.org/docs/current/static/sql-createtablesas.html>

# Temporary tables and views

Temporary tables are automatically dropped at the end of a session, or optionally at the end of the current transaction (see ON COMMIT below).

```
CREATE TEMP TABLE conference_and_workshop_papers AS  
  SELECT p.*  
  FROM papers p, venue v  
  WHERE p.venue = v.id AND v.type = 1;
```

Temporary views are automatically dropped at the end of the current session. If any of the tables referenced by the view are temporary, the view is created as a temporary view.

```
CREATE TEMP VIEW books_and_thesis AS  
  SELECT p.*  
  FROM papers p, venue v  
  WHERE p.venue = v.id AND v.type = 3;
```

For more information:

- <https://www.postgresql.org/docs/current/static/sql-createtable.html>
- <https://www.postgresql.org/docs/current/static/sql-createview.html>



# Full-text search

Full Text Search provides the capability to identify natural-language documents that satisfy a query, and optionally to sort them by relevance to the query.

Differently from textual data types operators such as LIKE, full text indexing allows documents to be preprocessed and an index saved for later rapid searching. Preprocessing includes:

- Parsing documents into tokens.
- Converting tokens into lexemes.
- Storing preprocessed documents optimized for searching.

It's possible to perform full text search without an index.

For more information:


- <https://www.postgresql.org/docs/current/static/textsearch-intro.html>

# Full-text search

```
SELECT id, name
FROM papers
WHERE to_tsvector('english', name) @@ to_tsquery('english', 'solve & problem')
LIMIT 7;
```

id	name
762788	Declaratively <b>solving</b> tricky Google Code Jam <b>problems</b> with Prolog-based ECLiPSe CLP system.
763157	Quantum Algorithms for many-to-one Functions to <b>Solve</b> the Regulator and the Principal Ideal <b>Problem</b>
763402	A Critique of " <b>Solving</b> the P/NP <b>Problem</b> Under Intrinsic Uncertainty".
763780	<b>Solving</b> the Parity <b>Problem</b> with Rule 60 in Array Size of the Power of Two.
764236	Towards <b>Solving</b> the Inverse Protein Folding <b>Problem</b>
764710	Induction of High-level Behaviors from <b>Problem-solving</b> Traces using Machine Learning Tools
764822	Abstract flows over time: A first step towards <b>solving</b> dynamic packing <b>problems</b>

For more information:

 <https://www.postgresql.org/docs/current/static/textsearch-controls.html>

# Full-text search

```
SELECT to_tsquery('english', 'the & solve & problem');
```

```
to_tsquery
```

---

```
'solv ' & 'problem '
```

```
SELECT id, name
```

```
FROM papers
```

```
WHERE to_tsvector('english', name) @@ to_tsquery('english', 'the & solve & problem')
```

```
LIMIT 4;
```

id	name
767312	A priori estimation of a time step for numerically solving parabolic problems.
767510	A novel approach of solving the CNF-SAT problem.
767645	Solving reviewer assignment problem in software peer review: An approach based on preference matrix and asymmetric TSP model.
767967	Ensuring Trust in One Time Exchanges: Solving the QoS Problem

For more information:

• <https://www.postgresql.org/docs/current/static/textsearch-controls.html>

# Full-text indexes

GIN (Generalized Inverted Index)-based index, contain an index entry for each word (lexeme), with a compressed list of matching locations.

The column must be of tsvector type.

```
CREATE INDEX ON papers USING GIN (to_tsvector('english', name));
```

```
CREATE INDEX ON venue USING GIN (to_tsvector('english', name || ' ' || school));
```

For more information:

- <https://www.postgresql.org/docs/current/static/textsearch-indexes.html>
- <https://www.postgresql.org/docs/current/static/textsearch-tables.html#textsearch-tables-index>

# PostgreSQL functions

```
SELECT replace('abcdefabcdef', 'cd', 'XX');
```

replace

---

```
abXXefabXXef  
(1 row)
```

```
SELECT unnest(ARRAY[1,2]);
```

unnest

---

```
1  
2  
(2 rows)
```

For more information:

- <https://www.postgresql.org/docs/current/static/functions-string.html>
- <https://www.postgresql.org/docs/current/static/functions-array.html>

# PostgreSQL functions (cont)

```
SELECT plainto_tsquery('english', 'The Fat Rats');
```

```
plainto_tsquery
```

---

```
'fat ' & 'rat '  
(1 row)
```

```
SELECT * FROM generate_series(2,4);
```

```
generate_series
```

---

```
2  
3  
4
```

```
(3 rows)
```

For more information:

- <https://www.postgresql.org/docs/current/static/functions-textsearch.html>
- <https://www.postgresql.org/docs/current/static/functions-srf.html>

# Table functions: crosstab

Produces a "pivot table" (that is, multiple rows) with the value columns specified by a second query.

crosstab function is part of a PostgreSQL extension called tablefunc. To call the crosstab function, you must first enable the tablefunc extension by executing the following SQL command: **CREATE extension tablefunc;**

```
SELECT *
FROM crosstab(
    'SELECT v.name, v.year, COUNT(p.id) AS papers
    FROM venue v, papers p
    WHERE v.id = p.venue
    GROUP BY v.name, v.year
    ORDER BY 1, 2',
    'SELECT DISTINCT year
    FROM venue
    ORDER BY 1 DESC
    LIMIT 6')
AS final_result(venue TEXT, "2016" INTEGER, "2015" INTEGER, "2014" IN-
TEGER, "2013" INTEGER, "2012" INTEGER, "2011" INTEGER);
```

For more information:

● <https://www.postgresql.org/docs/current/static/tablefunc.html>

# Table functions: crosstab

The first query parameter must be compliant with the following restrictions:

- The SELECT must return 3 columns.
- The first column in the SELECT will be the identifier of every row in the pivot table or final result. In our example, this is the venue's name.
- The second column in the SELECT represents the categories in the pivot table. In our example, these categories are the venue years.
- The third column in the SELECT represents the value to be assigned to each cell of the pivot table. In our example, these are the number of papers.

```
SELECT v.name, v.year, COUNT(p.id) AS papers
FROM venue v, papers p
WHERE v.id = p.venue
GROUP BY v.name, v.year
ORDER BY 1, 2
```

For more information:

• <https://www.postgresql.org/docs/current/static/tablefunc.html>



# Table functions: crosstab

The second query parameter represents the complete list of categories, in this case the last 6 years:

```
SELECT DISTINCT year  
FROM venue  
ORDER BY 1 DESC  
LIMIT 6
```

The crosstab function is invoked in the SELECT statement's FROM clause. We must define the names of the columns and data types that will go into the final result.

```
venue TEXT, "2016" INTEGER, "2015" INTEGER, "2014" INTEGER, "2013"  
INTEGER, "2012" INTEGER, "2011" INTEGER
```

For more information:

● <https://www.postgresql.org/docs/current/static/tablefunc.html>

# Table expressions: GROUPING SETS

Enables complex grouping operations.

The data selected by the FROM and WHERE clauses is grouped separately by each specified grouping set, aggregates computed for each group just as for simple GROUP BY clauses, and then the results returned.

```
SELECT v.name, v.year, COUNT(p.id) AS papers  
FROM venue v, papers p WHERE v.id = p.venue  
GROUP BY GROUPING SETS ((v.name), (v.year), ());
```

Each sublist of GROUPING SETS may specify zero or more columns or expressions and is interpreted the same way as though it were directly in the GROUP BY clause.

An empty grouping set means that all rows are aggregated down to a single group.

For more information:

- <https://www.postgresql.org/docs/current/static/queries-table-expressions.html#queries-grouping-sets>

# Table expressions: ROLLUP

A shorthand notation for a common type of grouping set.

**ROLLUP** (e1, e2, e3, ...)

represents the given list of expressions and all prefixes of the list including the empty list. Thus it is equivalent to:

```
GROUPING SETS (  
    (e1, e2, e3, ...),  
    ...  
    (e1, e2),  
    (e1),  
    ()  
)
```

Commonly used for analysis over hierarchical data; e.g. total population by city, state, and country.

For more information:

- <https://www.postgresql.org/docs/current/static/queries-table-expressions.html#queries-grouping-sets>

# Table expressions: CUBE

A shorthand notation for a common type of grouping set.

**CUBE**(e1, e2, e3)

represents the given list and all of its possible subsets (i.e. the power set).  
Thus it is equivalent to:

```
GROUPING SETS (  
    ( e1, e2, e3),  
    ( e1, e2 ),  
    ( e1, e3),  
    ( e1 ),  
    ( e2, e3),  
    ( e2 ),  
    ( e3 ),  
    ( )  
)
```

For more information:

- <https://www.postgresql.org/docs/current/static/queries-table-expressions.html#queries-grouping-sets>